

run

October 7, 2023

# 1 Using Machine Learning to predict Financial Crises

**Author:** [Chris Reimann](#) **Date created:** 2023/07/28 **Last modified:** 2023/10/07 **Description:** This notebook runs the experiments of the paper “Using Machine Learning to predict Financial Crises: An Evaluation of different Learning Algorithms for Early Warning Models”.

## 1.1 Setup

Load the latest code from the [repository](#) and import it into the virtual machine.

```
[2]: from prepareData import Data
      from doExperiment import Experiment
      import pandas as pd
```

## 1.2 Load Data

Construct datasets using specified indicator variables from the [MacroHistory](#) database. Nominal values of local currencies are transformed to GDP-ratios, while growth rates are computed for percentages and index values. For additional details see Chapter 2.3 of the paper.

```
[4]: # Define Indicator Sets
iv_macro = ["rconsbarro", "iy", "money", "xrusd", "cpi", "ca"]
iv_credit = ["tloans", "debtServ", "yieldCurve", "ltd", "debtgdp",
            ↪ "globaltloans", "globalyieldCurve"]
iv_ca = iv_credit + ["hpnom"]
iv_all = iv_macro + iv_ca

# Construct Datasets
df_macro = Data(indicators = iv_macro, crisisData = "MacroHistory").
            ↪getReady("Macro")
df_credit = Data(indicators = iv_credit, crisisData = "MacroHistory").
            ↪getReady("Credit")
df_ca = Data(indicators = iv_ca, crisisData = "MacroHistory").getReady("Credit",
            ↪ & Asset")
df_all = Data(indicators = iv_all, crisisData = "MacroHistory").getReady("All")
```

Macro: The final dataset contains 1591 observations with 63 distinct crisis events.

Credit: The final dataset contains 1373 observations with 60 distinct crisis

events.

Credit & Asset: The final dataset contains 1159 observations with 46 distinct crisis events.

All: The final dataset contains 1101 observations with 41 distinct crisis events.

### 1.3 Construct Experiments

Specify models to be tested. Available models are Logit, KNeighbors, RandomForest, ExtraTrees, SVM and NeuralNet.

```
[6]: # Specify Models to be tested
models = ["Logit", "KNeighbors", "RandomForest", "ExtraTrees", "SVM",
↪ "NeuralNet"]

# Define Experiments for all Indicator Sets.
ex_macroIS = Experiment(df_macro, models, "InSample")
ex_creditIS = Experiment(df_credit, models, "InSample")
ex_caIS = Experiment(df_ca, models, "InSample")
ex_allIS = Experiment(df_all, models, "InSample")

ex_macro = Experiment(df_macro, models, "CrossVal")
ex_credit = Experiment(df_credit, models, "CrossVal")
ex_ca = Experiment(df_ca, models, "CrossVal")
ex_all = Experiment(df_all, models, "CrossVal")
```

### 1.4 Run Experiments

#### 1.4.1 In-Sample

Compute ROC values for all models trained and tested on the whole dataset.

```
[4]: ex_macroIS.run(disableTqdm = True)
ex_creditIS.run(disableTqdm = True)
ex_caIS.run(disableTqdm = True)
ex_allIS.run(disableTqdm = True)

resIS = pd.concat([ex_macroIS.auc, ex_creditIS.auc, ex_caIS.auc, ex_allIS.auc],
↪ axis = 0)
resIS
```

```
[4]:
```

	Set	Model	AUC
0	Macro	KNeighbors	1.000000
1	Macro	RandomForest	0.975804
2	Macro	ExtraTrees	0.825952
3	Macro	SVM	0.773601
4	Macro	Logit	0.709806
5	Macro	NeuralNet	0.707416
6	Macro	Random Assignment	0.500000

0	Credit	KNeighbors	1.000000
1	Credit	RandomForest	0.983538
2	Credit	ExtraTrees	0.934343
3	Credit	SVM	0.918643
4	Credit	NeuralNet	0.851912
5	Credit	Logit	0.790217
6	Credit	Random Assignment	0.500000
0	Credit & Asset	KNeighbors	1.000000
1	Credit & Asset	RandomForest	0.992726
2	Credit & Asset	ExtraTrees	0.955384
3	Credit & Asset	SVM	0.943323
4	Credit & Asset	NeuralNet	0.864574
5	Credit & Asset	Logit	0.808380
6	Credit & Asset	Random Assignment	0.500000
0	All	KNeighbors	1.000000
1	All	RandomForest	0.993947
2	All	ExtraTrees	0.963143
3	All	SVM	0.949269
4	All	NeuralNet	0.927175
5	All	Logit	0.858661
6	All	Random Assignment	0.500000

### 1.4.2 Out-of-Sample

Run experiments in cross-validation setting.

```
[5]: n = 100 # Specify number of cross-validation iterations

ex_macro.run(n)
ex_credit.run(n)
ex_ca.run(n)
ex_all.run(n)

resCrossVal = pd.concat([ex_macro.auc, ex_credit.auc, ex_ca.auc, ex_all.auc])
resCrossVal
```

```
Macro: Random Assignment:  0%|          | 0/100 [00:00<?, ?it/s]
Macro: Logit:             0%|          | 0/100 [00:00<?, ?it/s]
Macro: KNeighbors:        0%|          | 0/100 [00:00<?, ?it/s]
Macro: RandomForest:      0%|          | 0/100 [00:00<?, ?it/s]
Macro: ExtraTrees:        0%|          | 0/100 [00:00<?, ?it/s]
Macro: SVM:               0%|          | 0/100 [00:00<?, ?it/s]
Macro: NeuralNet:         0%|          | 0/100 [00:00<?, ?it/s]
Credit: Random Assignment: 0%|          | 0/100 [00:00<?, ?it/s]
```

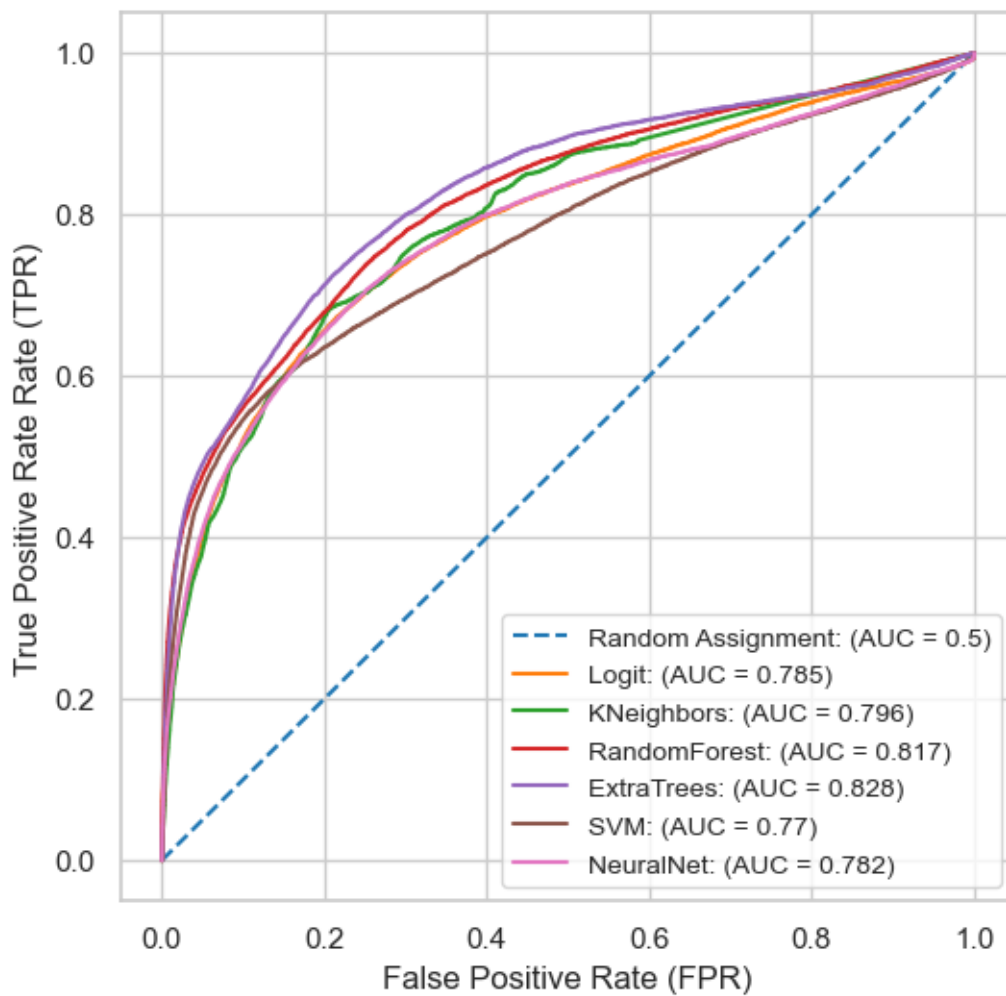
Credit: Logit: 0%| | 0/100 [00:00<?, ?it/s]  
 Credit: KNeighbors: 0%| | 0/100 [00:00<?, ?it/s]  
 Credit: RandomForest: 0%| | 0/100 [00:00<?, ?it/s]  
 Credit: ExtraTrees: 0%| | 0/100 [00:00<?, ?it/s]  
 Credit: SVM: 0%| | 0/100 [00:00<?, ?it/s]  
 Credit: NeuralNet: 0%| | 0/100 [00:00<?, ?it/s]  
 Credit & Asset: Random Assignment: 0%| | 0/100 [00:00<?, ?it/s]  
 Credit & Asset: Logit: 0%| | 0/100 [00:00<?, ?it/s]  
 Credit & Asset: KNeighbors: 0%| | 0/100 [00:00<?, ?it/s]  
 Credit & Asset: RandomForest: 0%| | 0/100 [00:00<?, ?it/s]  
 Credit & Asset: ExtraTrees: 0%| | 0/100 [00:00<?, ?it/s]  
 Credit & Asset: SVM: 0%| | 0/100 [00:00<?, ?it/s]  
 Credit & Asset: NeuralNet: 0%| | 0/100 [00:00<?, ?it/s]  
 All: Random Assignment: 0%| | 0/100 [00:00<?, ?it/s]  
 All: Logit: 0%| | 0/100 [00:00<?, ?it/s]  
 All: KNeighbors: 0%| | 0/100 [00:00<?, ?it/s]  
 All: RandomForest: 0%| | 0/100 [00:00<?, ?it/s]  
 All: ExtraTrees: 0%| | 0/100 [00:00<?, ?it/s]  
 All: SVM: 0%| | 0/100 [00:00<?, ?it/s]  
 All: NeuralNet: 0%| | 0/100 [00:00<?, ?it/s]

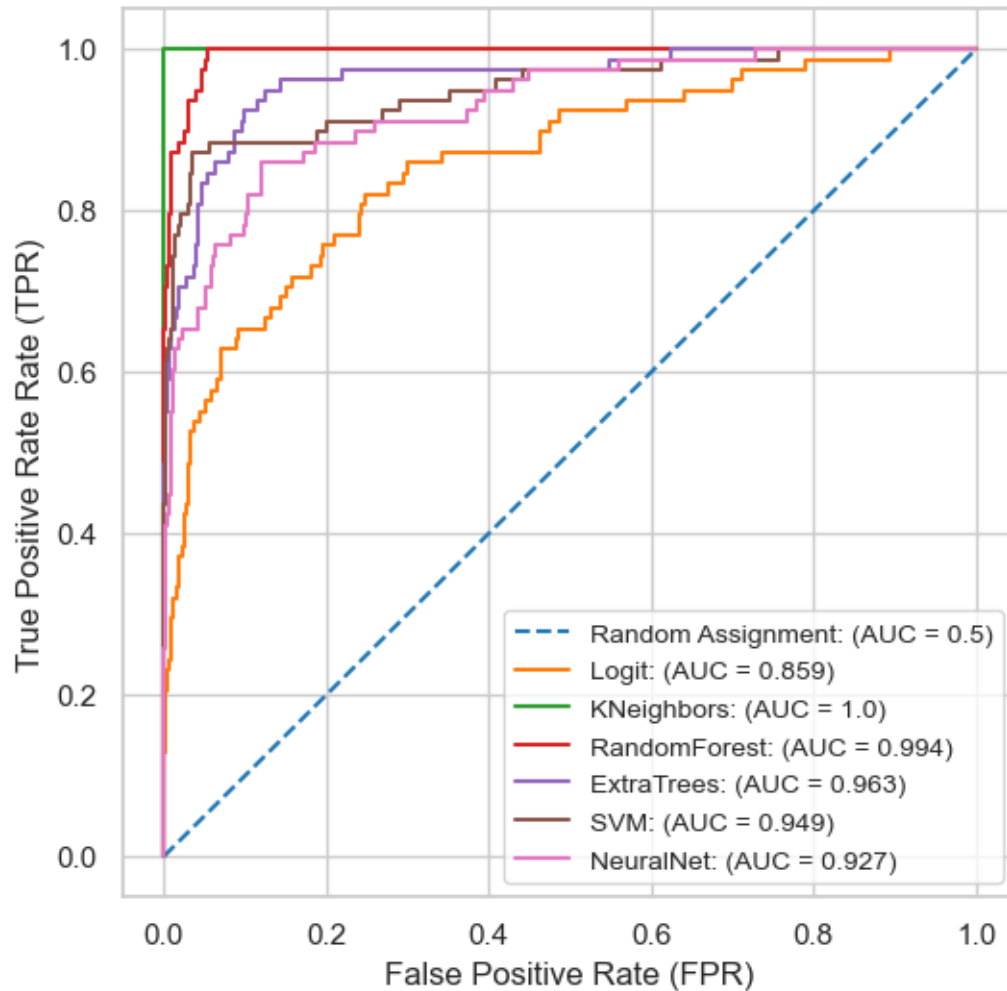
[5]:

	Set	Model	AUC
0	Macro	RandomForest	0.704588
1	Macro	Logit	0.679660
2	Macro	NeuralNet	0.663692
3	Macro	ExtraTrees	0.656247
4	Macro	KNeighbors	0.603511
5	Macro	SVM	0.516354
6	Macro	Random Assignment	0.500000
0	Credit	ExtraTrees	0.825699
1	Credit	RandomForest	0.824423
2	Credit	KNeighbors	0.799671
3	Credit	NeuralNet	0.772014
4	Credit	Logit	0.755909
5	Credit	SVM	0.751870
6	Credit	Random Assignment	0.500000
0	Credit & Asset	ExtraTrees	0.839913
1	Credit & Asset	RandomForest	0.832740

2	Credit & Asset	KNeighbors	0.798748
3	Credit & Asset	NeuralNet	0.772401
4	Credit & Asset	Logit	0.763955
5	Credit & Asset	SVM	0.751551
6	Credit & Asset	Random Assignment	0.500000
0	All	ExtraTrees	0.828041
1	All	RandomForest	0.816599
2	All	KNeighbors	0.795852
3	All	Logit	0.785161
4	All	NeuralNet	0.781879
5	All	SVM	0.769958
6	All	Random Assignment	0.500000

```
[6]: ex_all.rocGraph()  
ex_allIS.rocGraph()
```

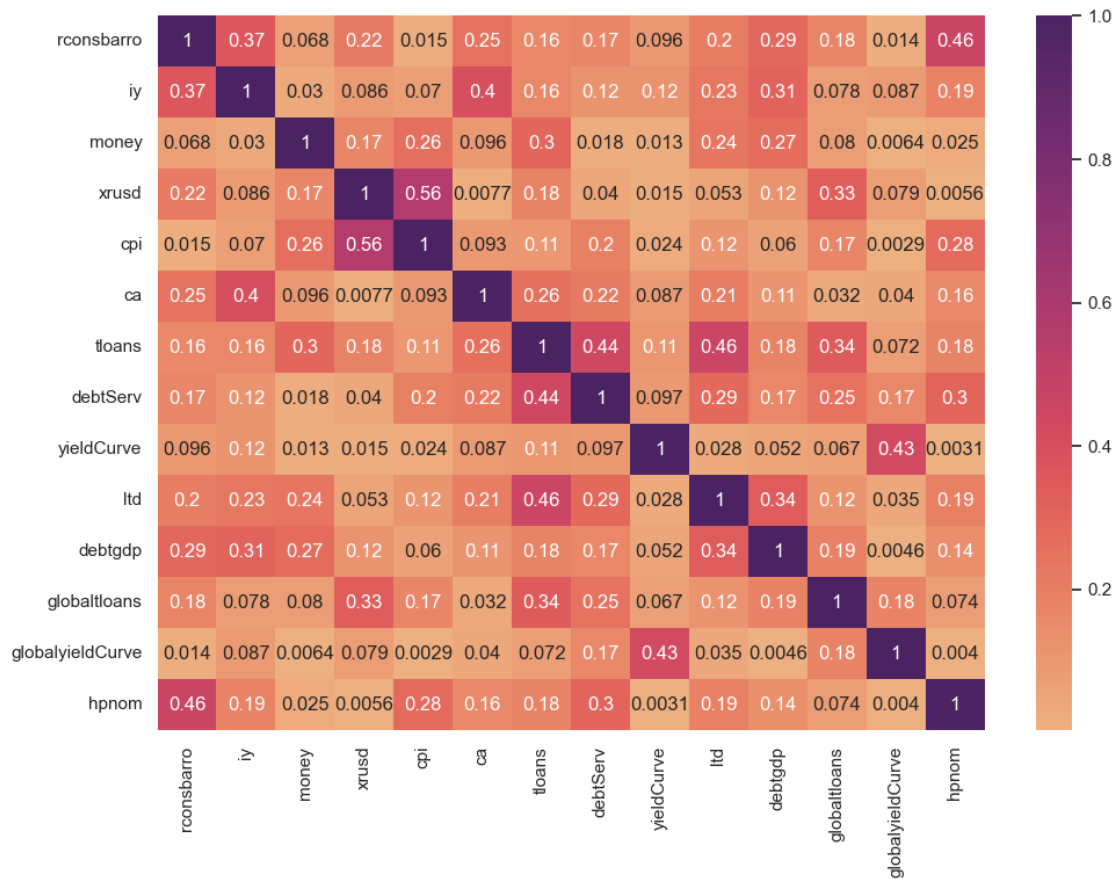




## 1.5 Explainability

Compute logistic regression coefficients and AEL for best performing black-blox model (Random Forest / Extremely Randomized Trees).

```
[7]: df_all.correlationMatrix()
```



```
[8]: df_all.vif()
```

```
[8]:
```

	variable	VIF
0	rconsbarro	3.139997
13	hpnom	2.501668
6	tloans	2.482708
12	globalyieldCurve	2.288070
4	cpi	2.273961
9	ltd	1.839108
2	money	1.764422
11	globaltloans	1.743562
3	xrusd	1.732725
8	yieldCurve	1.540389
1	iy	1.459956
7	debtServ	1.444777
10	debtgdp	1.355034
5	ca	1.302359

```
[9]: ex_all.logitCoef()
```

Optimization terminated successfully.

Current function value: 0.179190

Iterations 8

const is significant at 1%  
rconsbarro is significant at 1%  
money is significant at 5%  
xrusd is significant at 5%  
cpi is significant at 1%  
debtServ is significant at 5%  
yieldCurve is significant at 1%  
ltd is significant at 1%  
globaltloans is significant at 1%  
globalyieldCurve is significant at 1%  
hpnom is significant at 10%

[9]:	Coef.	Std.Err.	z	P> z	[0.025	\
const	-3.697644	0.228878	-16.155525	1.038316e-58	-4.146237	
rconsbarro	-0.629918	0.191872	-3.283011	1.027045e-03	-1.005980	
iy	-0.051658	0.169415	-0.304922	7.604257e-01	-0.383706	
money	0.389246	0.174602	2.229332	2.579181e-02	0.047032	
xrusd	-0.410281	0.191527	-2.142155	3.218099e-02	-0.785668	
cpi	0.748978	0.171188	4.375175	1.213349e-05	0.413456	
ca	-0.258403	0.158231	-1.633073	1.024535e-01	-0.568530	
tloans	0.223606	0.180472	1.239010	2.153419e-01	-0.130112	
debtServ	-0.401652	0.176307	-2.278137	2.271840e-02	-0.747208	
yieldCurve	-0.787079	0.176864	-4.450203	8.578928e-06	-1.133725	
ltd	0.531080	0.173788	3.055914	2.243756e-03	0.190462	
debtgdp	-0.170815	0.192149	-0.888970	3.740192e-01	-0.547421	
globaltloans	0.756404	0.172335	4.389149	1.137953e-05	0.418633	
globalyieldCurve	-0.583369	0.160648	-3.631355	2.819367e-04	-0.898233	
hpnom	0.236903	0.121942	1.942752	5.204617e-02	-0.002099	

0.975]

const	-3.249052
rconsbarro	-0.253856
iy	0.280389
money	0.731459
xrusd	-0.034895
cpi	1.084501
ca	0.051724
tloans	0.577324
debtServ	-0.056096
yieldCurve	-0.440433
ltd	0.871697
debtgdp	0.205791
globaltloans	1.094174
globalyieldCurve	-0.268505



hpnom                    0.475905

```
[10]: ex_allIS.ALE(["Logit", "RandomForest", "ExtraTrees"], range(0,14))
```

```
LogisticRegression(max_iter=1000, penalty='none', random_state=1)
```

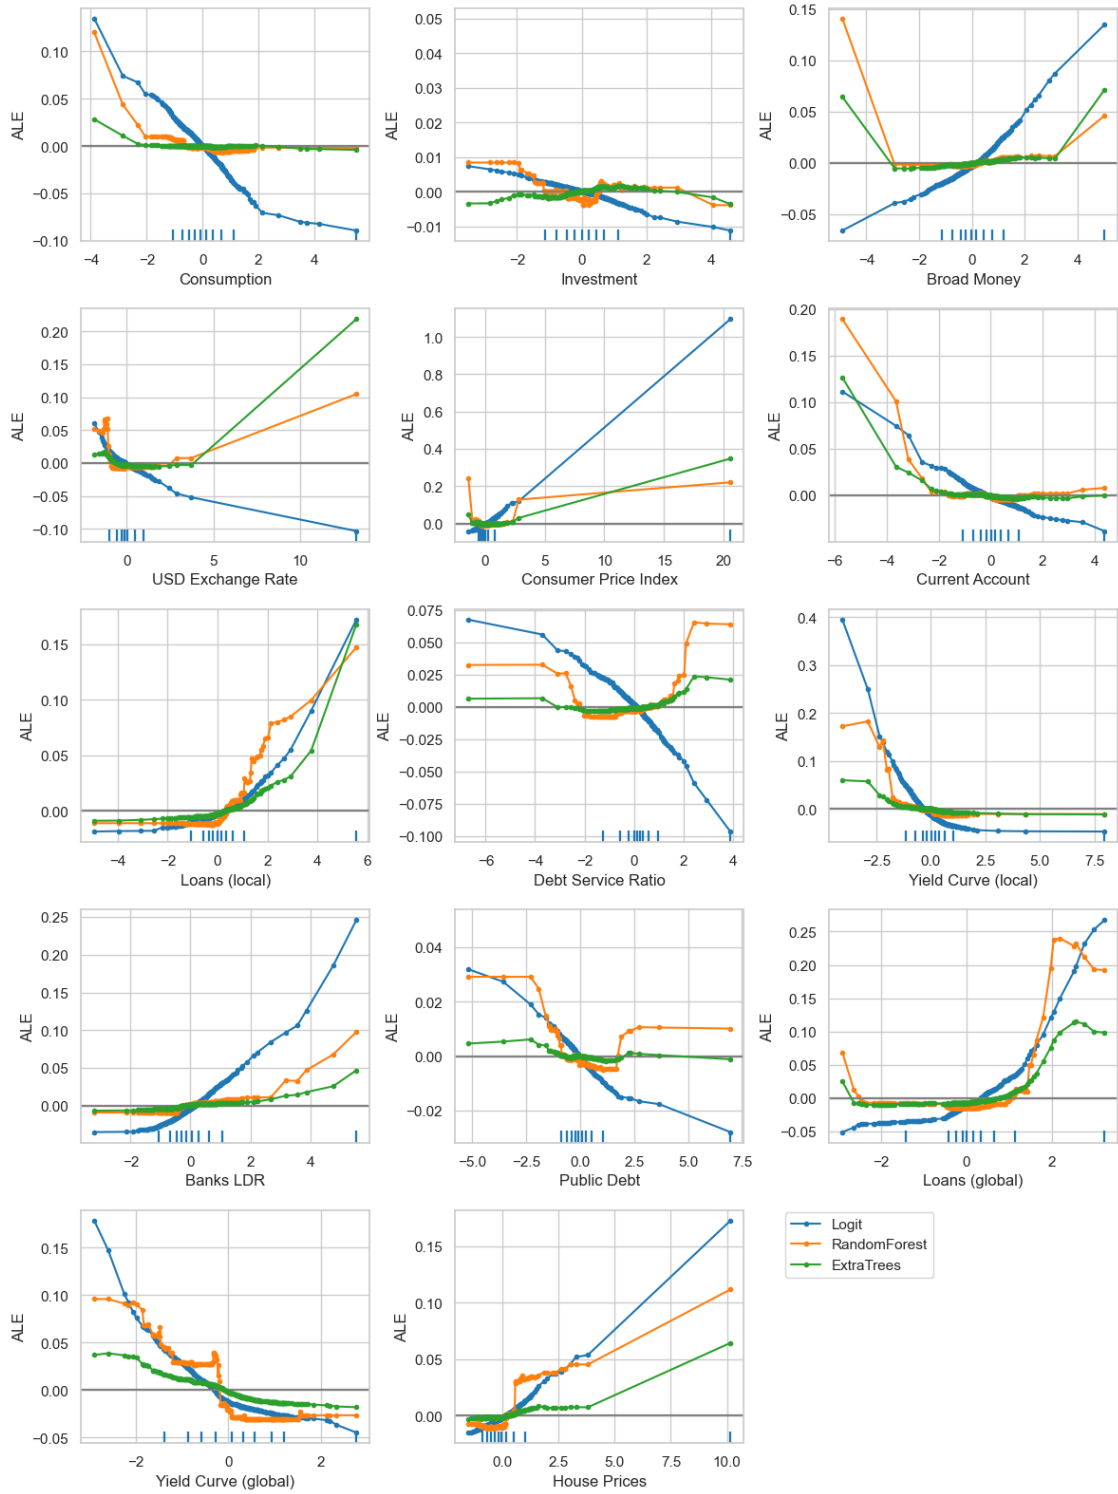
```
AUC: 0.8586610522094392
```

```
RandomForestClassifier(max_depth=6, n_estimators=1000, random_state=1)
```

```
AUC: 0.993946913301752
```

```
ExtraTreesClassifier(max_depth=6, n_estimators=1000, random_state=1)
```

```
AUC: 0.9631425921748502
```



## 1.6 Robustness Checks

```
[7]: # Robustness Check: ESRB Crisis Data
df_alt1 = Data(iv_all, crisisData = "ESRB").getReady("ESRB")
ex_alt1 = Experiment(df_alt1, models, "CrossVal")
ex_alt1.run(n = 100)
ex_alt1.auc
```

ESRB: The final dataset contains 426 observations with 22 distinct crisis events.

ESRB: Random Assignment: 0%| | 0/100 [00:00<?, ?it/s]

ESRB: Logit: 0%| | 0/100 [00:00<?, ?it/s]

ESRB: KNeighbors: 0%| | 0/100 [00:00<?, ?it/s]

ESRB: RandomForest: 0%| | 0/100 [00:00<?, ?it/s]

ESRB: ExtraTrees: 0%| | 0/100 [00:00<?, ?it/s]

ESRB: SVM: 0%| | 0/100 [00:00<?, ?it/s]

ESRB: NeuralNet: 0%| | 0/100 [00:00<?, ?it/s]

```
[7]:      Set      Model      AUC
0 ESRB      RandomForest 0.896315
1 ESRB      ExtraTrees  0.867543
2 ESRB      KNeighbors  0.798209
3 ESRB      Logit       0.770237
4 ESRB      SVM         0.764440
5 ESRB      NeuralNet   0.743398
6 ESRB      Random Assignment 0.500000
```

```
[8]: # Robustness Check: Laeven & Valencia Crisis Data
df_alt1 = Data(iv_all, crisisData = "LaevenValencia").getReady("LaevenValencia")
ex_alt1 = Experiment(df_alt1, models, "CrossVal")
ex_alt1.run(n = 100)
ex_alt1.auc
```

Yugoslavia, SFR not found in regex

LaevenValencia: The final dataset contains 674 observations with 18 distinct crisis events.

LaevenValencia: Random Assignment: 0%| | 0/100 [00:00<?, ?it/s]

LaevenValencia: Logit: 0%| | 0/100 [00:00<?, ?it/s]

LaevenValencia: KNeighbors: 0%| | 0/100 [00:00<?, ?it/s]

LaevenValencia: RandomForest: 0%| | 0/100 [00:00<?, ?it/s]

LaevenValencia: ExtraTrees: 0%| | 0/100 [00:00<?, ?it/s]

LaevenValencia: SVM: 0%| | 0/100 [00:00<?, ?it/s]

LaevenValencia: NeuralNet: 0%| | 0/100 [00:00<?, ?it/s]

```
[8]:
```

	Set	Model	AUC
0	LaevenValencia	KNeighbors	0.867680
1	LaevenValencia	ExtraTrees	0.867093
2	LaevenValencia	NeuralNet	0.863735
3	LaevenValencia	SVM	0.853616
4	LaevenValencia	RandomForest	0.851058
5	LaevenValencia	Logit	0.837715
6	LaevenValencia	Random Assignment	0.500000

```
[ ]:
```